

# Quick, Draw! Image Classification

Group 3

December 6, 2018

## 1 Introduction

Many recent industry innovations have included neural networks. An important research area within neural networks has been image classification. One of the best techniques to classify images is Convolutional Neural Networks (CNN). We have sought to better understand industry best practices by exploring different image classification techniques with neural networks. Our exploration starts by manually performing image classification and noting the level of human accuracy. We then use a fully connected neural network to perform image classification before implementing a CNN. We implemented options for tuning and debugging the neural network while implementing each method during our process. We were able to develop an intuition for image classification best practices using neural networks after our research.

Many supervised machine learning methods often differ from those employed when constructing neural networks. These supervised machine learning methods, such as logistic regression or decision tree classification, require feature engineering to successfully classify an image. Models requiring feature engineering rely on considerable effort by subject matter experts to achieve reasonable accuracy. Neural networks allow for an approach which does not require feature engineering to achieve similar or better accuracy. Industry engineering efforts regarding feature engineering can be shifted to neural networks in some cases such as image classification. Moreover, the availability of cheap computing resources sped up a shift towards the use of neural networks.

When considering modeling approaches, we examine both fully connected and convolutional neural networks. A fully connected neural network is known to be inefficient at classifying large images. These inefficiencies in a fully connected neural network partly arise because every neuron in the  $\ell - 1$  layer must be connected to the  $\ell$ th layer. This requires a series of matrix multiplications and additions across each layer  $\ell$ . Alternatively, a CNN allows for more efficient computations by: (1) efficient and automatic feature generation using local connectivity and parameter sharing of convolution operations, (2) dimensionality reduction using pooling layers. In this project, we attempt to understand these topics in depth from an empirical standpoint in relation to image classification using Google's "Quick, Draw!" dataset.

## 2 Methods

We generate a series of hypotheses to inform our analysis.

### 2.1 Hypotheses

We start off with hypothesis that we learnt from theoretical literature. We then perform several experiments to test the validity of these hypothesis and report out our findings. Some of our hypothesis is as follows:

- 1) Fully connected neural network will struggle to get reasonable accuracy. In contrast CNN's should perform better and faster than a Fully connected neural network.
- 2) Training a deep neural network with limited data will result in overfitting. Conversely, adding more data to a deep neural network is an effective way to manage variance.
- 3) Imbalance class prediction can be improved using data augmentation.
- 4) Increasing the number of categories will cause a significant drop in the performance of the neural network. Increasing the overall data points for the network to be trained on should help.

### 2.2 Data

Google originally collected the data from users by explicitly asking them to draw objects like forks, hotdogs, etc. The original dataset is 70 GB in size but Google also provides a simplified version of the dataset- consisting only of the final images. We used these '.npy' files to train all of our models. These numpy bitmap files consist of more than 100,000 rows each and 784 features, since they are basically  $28 \times 28$  images. Samples of these images are shown in Figure 1. The subset of this dataset sampled varies from experiment to experiment. More details can be found in the experiments mentioned below.



Figure 1: Sample images from Quick, Draw! dataset

## 3 Results

### 3.1 Hypothesis 1

Hypothesis 1: Fully connected neural network will struggle to get reasonable accuracy. In contrast CNN's should perform better and faster than a Fully connected neural network.

In order to perform this experiment we built a fully connected neural network using tensorflow. Our dataset sample selected for this experiment had 10 categories: fish, fork, hotdog, flamingo, airplane, alarm clock, baseball bat, bicycle, dolphin, elephant. We had 100,000 examples in total, with 10,000 example for each of the aforementioned categories. Our train:test split was chosen to be 80:20, an industry standard.

Our initial objective was to build a model that overfit our given dataset, with the idea that once our model is complex enough to overfit the training data, we could regularize it in order to achieve a more optimal fit.

After experimenting with various architectures, we narrowed down to a 3 layer model as specified in Figure 2.

$$Z_1 = \underset{25 \times 784}{W_1} X + \underset{25 \times 1}{b_1} \Rightarrow A_1 = \text{ReLU}(Z_1)$$

$$Z_2 = \underset{12 \times 25}{W_2} A_1 + \underset{12 \times 1}{b_2} \Rightarrow A_2 = \text{ReLU}(Z_2)$$

$$Z_3 = \underset{10 \times 12}{W_3} A_2 + \underset{10 \times 1}{b_3} \Rightarrow \text{output layer} = \text{SoftMax}(Z_3)$$

$$\underset{Z_1}{(25 \times 784 + 25 \times 1)} + \underset{Z_2}{(12 \times 25 + 12 \times 1)} + \underset{Z_3}{(10 \times 12 + 10 \times 1)} = 20,067 \text{ parameters}$$

Figure 2: Fully Connected Neural Network Architecture

The training set accuracy observed was 93%, while the test accuracy was at 73%, which clearly indicated an overfit in the model, as designed.

The model took 1500 epochs to converge taking around 45 minutes on a machine with 13 GB of RAM, Intel Xeon CPU 2.3 GHz (1 core, 2 threads), 1xTesla K80 GPU having 2496 CUDA cores with 12 GB GDDR5 VRAM. In order to manage variance, we attempted regularizing using dropouts. Unfortunately, that effort was futile. We then doubled the amount of data sampled to train the model, eventually helping us achieve training set accuracy of 86% and test set accuracy of 82%. Our learning rate was specified as 0.001.

While our initial hypothesis assumed that a regular fully connected network would struggle to achieve a reasonable accuracy, our experiment proved our hypothesis as invalid by achieving a commendable accuracy. Having said that, the model did take a longer to converge than

28 x 28 x 1 [Input Layer]	3 x 3 x 4 [ReLU] [Convolution]	2 x 2 x 8 [ReLU] [Convolution]	64 [Relu] [Dense]	10 [Softmax] [Output Layer]
------------------------------	--------------------------------------	--------------------------------------	-------------------------	-----------------------------------

Figure 3: CNN Architecture

28 x 28 x 1 [Input Layer]	3 x 3 x 4 [ReLU] [Convolution] [MaxPool - 2x2]	2 x 2 x 8 [ReLU] [Convolution] [MaxPool - 2x2]	64 [Relu] [Dense]	10 [Softmax] [Output Layer]
------------------------------	---	---	-------------------------	-----------------------------------

Figure 4: CNN Architecture with Reduced Parameters

expected.

As for the CNN, we started off with the same approach with attempting to overfit the model using convolution operations and a fully connected layer. Since our objective is to categorise 10 categories, softmax was chosen as the activation function for the final layer, while every other neuron had ReLU as the activation function owing to ReLU's faster convergence.

Loss Function: Categorical Cross Entropy [TO DO: Why?] Optimizer: [TODO: Why?]  
Number of Trainable Parameters: 320,890

The number of clearly being much greater than the number of observations used to train the model, we need to reduce the number of parameters vis-a-vis number of observations. From the available strategies to do so, we incorporated MaxPooling [TO DO: why?] post each convolution operation, in order to extract only the most significant features. As a consequence, this reduced the volume of the output, giving the model a smaller memory footprint as well as a smaller number of trainable parameters.

The observable training accuracy is 0.90, whereas the validation accuracy is 0.89. The model converged after 10 epochs, which took 3 minutes.

Some of the other unsuccessful strategies used to decrease the number of parameters and to regularize the network were strides and dropouts which sacrificed way too much in the way of accuracy and thus had to be dropped. Since most of our drawings had most of their information in the center of the image, padding, in theory wouldn't be of much help. However, actually implementing padding improved our accuracy by 1% at a cost of 6000 parameters.

## Conclusion

A slight (5%) improvement in the accuracy of a CNN when compared to a fully connected network, keeping the number of parameters constant, was observed. However, the CNN happened to be around 13 times faster in practice. This delta in accuracy and speed is likely to be amplified as we scale the model up with respect to model complexity or data.

## **3.2 Hypothesis 2**

Hypothesis 2: Training a deep neural network with limited data will result in overfitting. Conversely, adding more data to a deep neural network is an effective way to manage variance.

With us reaching an optimal model for the dataset, we decided to study the effect of reduction in the sampled dataset and its relationship with overfitting.

## **3.3 Hypothesis 3**

We tried to show the importance of having equal partitions of data to train a machine learning model. Imbalanced classes can make a model biased toward the majority class.

## **4 Discussion**